

1 Homework

1.1 Import

```
[3]: %matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import style
from sklearn.datasets import fetch_openml, make_classification
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from mpl_toolkits.mplot3d import Axes3D
```

1.2 1.7

```
[4]: def check_bmi(bmi):
    if bmi < 18.5:
        return 'Underweight'
    elif bmi < 25:
        return 'Normal weight'
    elif bmi < 30:
        return 'Overweight'
```

```
[5]: names = np.array(['Ann', 'Joe', 'Mark'])
heights = np.array([1.5, 1.78, 1.6])
weights = np.array([65, 46, 59])

bmi = weights / heights ** 2
bmi
```

```
[5]: array([28.88888889, 14.51836889, 23.046875  ])
```

```
[6]: df = pd.DataFrame({'Name': names, 'Height': heights, 'Weight': weights, 'BMI':
    ↪bmi})
df
```

```
[6]:
```

	Name	Height	Weight	BMI
0	Ann	1.50	65	28.888889
1	Joe	1.78	46	14.518369
2	Mark	1.60	59	23.046875

```
[7]: classify = np.vectorize(check_bmi)
classify(bmi)
```

```
[7]: array(['Overweight', 'Underweight', 'Normal weight'], dtype='<U13')
```

```
[8]: df2 = pd.DataFrame({
    'Name': names,
    'Height': heights,
    'Weight': weights,
    'BMI': bmi,
    'Classify': classify(bmi)})
df2
```

```
[8]:      Name  Height  Weight      BMI      Classify
0   Ann    1.50     65  28.888889   Overweight
1   Joe    1.78     46  14.518369  Underweight
2  Mark    1.60     59  23.046875  Normal weight
```

1.2.1 Data from group

```
[9]: data = pd.read_csv('data/no1_7.csv')
data
```

```
[9]:      name  height  weight
0   anhnt    1.66     72
1  vphuong    1.78     65
2      vu    1.68     60
3     nam    1.69     65
4  dphuong    1.67     60
```

```
[10]: names = data['name'].values
heights = data['height'].values
weights = data['weight'].values
```

```
[11]: bmi = weights / heights ** 2
bmi
```

```
[11]: array([26.12861083, 20.51508648, 21.2585034 , 22.75830678, 21.51385851])
```

```
[12]: df3 = pd.DataFrame({
    'Name': names,
    'Height': heights,
    'Weight': weights,
    'BMI': bmi,
    'Classify': classify(bmi)})
df3
```

```
[12]:      Name  Height  Weight      BMI      Classify
0   anhnt    1.66     72  26.128611   Overweight
1  vphuong    1.78     65  20.515086  Normal weight
2      vu    1.68     60  21.258503  Normal weight
3     nam    1.69     65  22.758307  Normal weight
4  dphuong    1.67     60  21.513859  Normal weight
```


1.3 1.8

Plotting multiple Lines in the same chart

```
[13]: style.use('ggplot')
```

```
[14]: plt.plot(  
    [1,2,3,4,5,6,7,8,9,10],  
    [2,4.5,1,2,3.5,2,1,2,3,2]  
)  
  
plt.plot(  
    [1,2,3,4,5,6,7,8,9,10],  
    [3,4,2,5,2,4,2.5,4,3.5,3]  
)  
  
plt.title('Results')  
plt.xlabel('Semester')  
plt.ylabel('Grade')
```

```
[14]: Text(0, 0.5, 'Grade')
```




Adding a Legend

```
[15]: plt.plot(
      [1,2,3,4,5,6,7,8,9,10],
      [2,4.5,1,2,3.5,2,1,2,3,2],
      label="Jim"
    )
    plt.plot(
      [1,2,3,4,5,6,7,8,9,10],
      [3,4,2,5,2,4,2.5,4,3.5,3],
      label="Tom"
    )

    plt.title('Results')
    plt.xlabel('Semester')
    plt.ylabel('Grade')
    plt.legend()
```

```
[15]: <matplotlib.legend.Legend at 0x1bfce2274c0>
```




Plotting bar charts

```
[16]: plt.bar(
        [1,2,3,4,5,6,7,8,9,10],
        [2,4.5,1,2,3.5,2,1,2,3,2],
        label = "Jim",
        color = "m",           # m for magenta
        align = "center"
    )

    plt.title("Results")
    plt.xlabel("Semester")
    plt.ylabel("Grade")

    plt.legend()
    plt.grid(True, color="y")
```



1.3.1 Data from team

```
[17]: df = pd.read_csv('./data/no1_8.csv')
df
```

```
[17]:
```

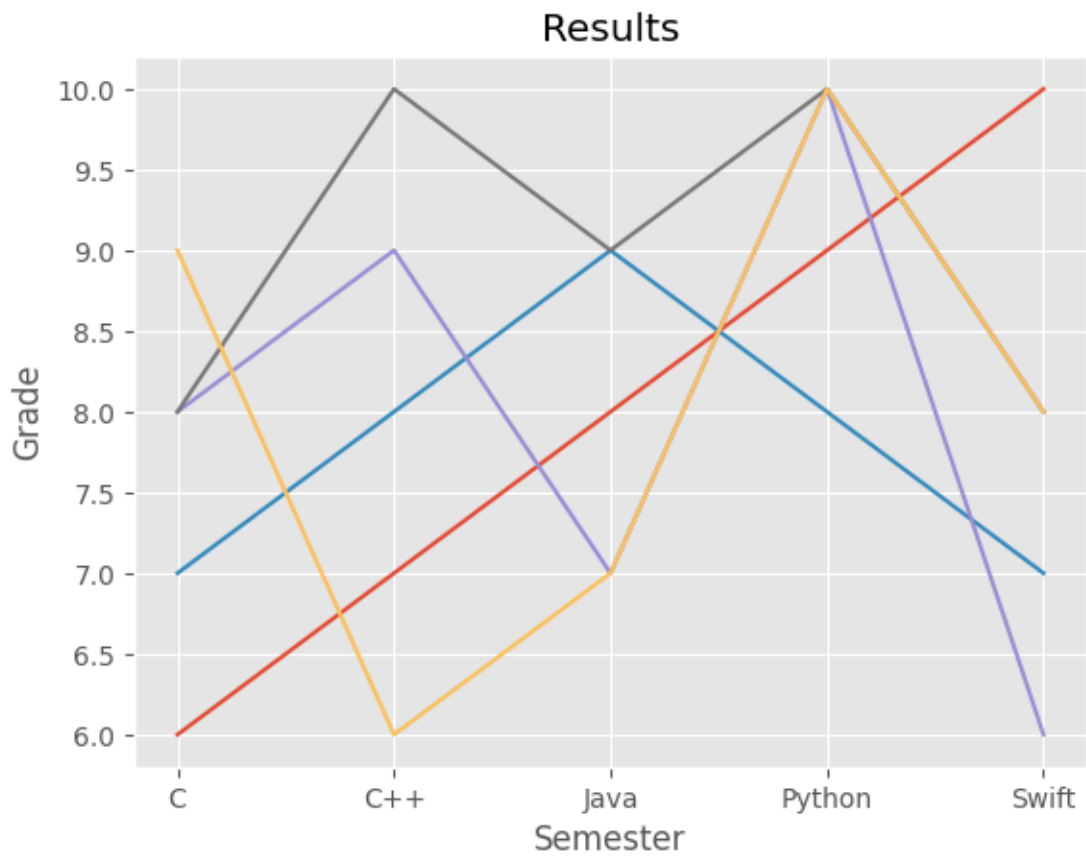
	Subject	tienanh	vphuong	vu	nam	dphuong
0	C	6	7	8	8	9
1	C++	7	8	9	10	6
2	Java	8	9	7	9	7
3	Python	9	8	10	10	10
4	Swift	10	7	6	8	8

```
[18]: subjects = df["Subject"]
tienanh_grade = df["tienanh"]
vphuong_grade = df["vphuong"]
vu_grade = df["vu"]
nam_grade = df["nam"]
dphuong_grade = df["dphuong"]
```

```
[19]: plt.plot(
    subjects,
    tienanh_grade,
    label="tienanh",
)
plt.plot(
    subjects,
    vphuong_grade,
    label="vphuong",
)
plt.plot(
    subjects,
    vu_grade,
    label="vu",
)
plt.plot(
    subjects,
    nam_grade,
    label="nam",
)
plt.plot(
    subjects,
    dphuong_grade,
    label="dphuong",
)

plt.title('Results')
plt.xlabel('Semester')
plt.ylabel('Grade')
```

```
[19]: Text(0, 0.5, 'Grade')
```




```
[20]: style.use('ggplot')

subjects = subjects.tolist()
subjects

plt.bar(
    subjects,
    tienanh_grade,
    label="tienanh",
    color = "m",
    align = "center"
)

plt.title('Results')
plt.xlabel('Semester')
plt.ylabel('Grade')
```

```
plt.legend()
plt.grid(True, color="y")
```



```
[21]: plt.plot(
    subjects,
    tienanh_grade,
    label="tienanh",
)
plt.plot(
    subjects,
    vphuong_grade,
    label="vphuong",
)
plt.plot(
    subjects,
    vu_grade,
    label="vu",
)
plt.plot(
```




```

subjects,
nam_grade,
label="nam",
)
plt.plot(
subjects,
dphuong_grade,
label="dphuong",
)

plt.xlabel("Subject")
plt.ylabel("Grade")
plt.title("Grade of Students")
plt.legend()

```

[21]: <matplotlib.legend.Legend at 0x1bfcf2cbeb0>



1.4 1.9

```
[22]: df = pd.read_csv('data/no1_9.csv')
df
```

```
[22]:
```

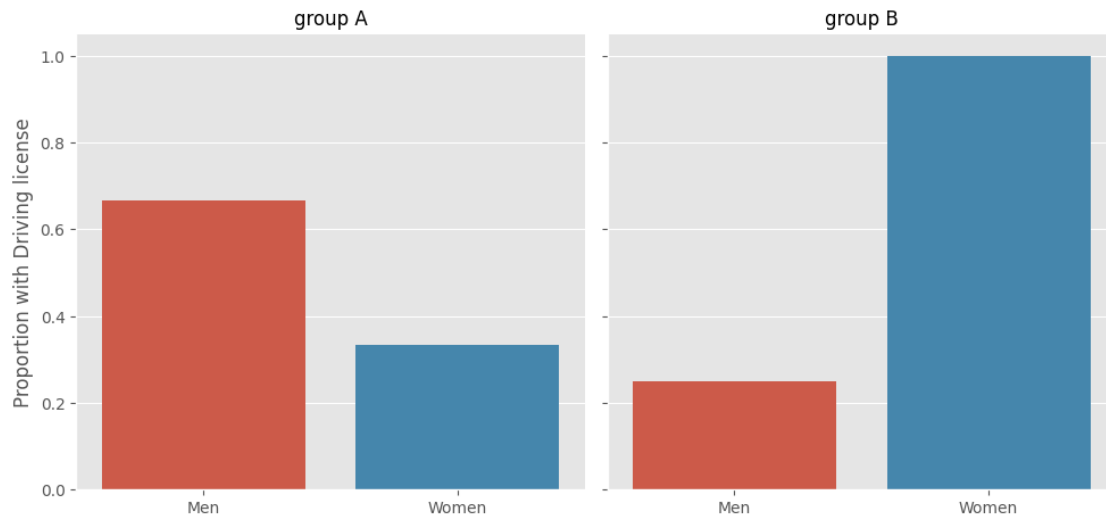
	gender	group	license
0	men	A	1
1	men	A	0
2	men	A	1
3	women	A	1
4	women	A	0
5	women	A	0
6	men	B	0
7	men	B	0
8	men	B	0
9	men	B	1
10	women	B	1
11	women	B	1
12	women	B	1
13	women	B	1

```
[23]: g = sns.catplot(x="gender", y="license", col="group",
                    data = df, kind="bar", errorbar=None, aspect=1.0)

#--- set the labels ---
g.set_axis_labels("", "Proportion with Driving license")
g.set_xticklabels(["Men", "Women"])
g.set_titles("{col_var} {col_name}")

#--- show plot ---
plt.show()
```

```
c:\Users\tien2\miniconda3\lib\site-packages\seaborn\axisgrid.py:118:
UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)
```



1.5 1.10

```
[24]: df = pd.read_csv('data/no1_10.csv')
df
```

```
[24]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class \
0	0	3	male	22.0	1	0	7.2500	S	Third
1	1	1	female	38.0	1	0	71.2833	C	First
2	1	3	female	26.0	0	0	7.9250	S	Third
3	1	1	female	35.0	1	0	53.1000	S	First
4	0	3	male	35.0	0	0	8.0500	S	Third
..
886	0	2	male	27.0	0	0	13.0000	S	Second
887	1	1	female	19.0	0	0	30.0000	S	First
888	0	3	female	NaN	1	2	23.4500	S	Third
889	1	1	male	26.0	0	0	30.0000	C	First
890	0	3	male	32.0	0	0	7.7500	Q	Third

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True
..
886	man	True	NaN	Southampton	no	True
887	woman	False	B	Southampton	yes	True
888	woman	False	NaN	Southampton	no	False
889	man	True	C	Cherbourg	yes	True

```
890    man          True  NaN  Queenstown    no    True
```


```
[891 rows x 15 columns]
```

```
[25]: g = sns.catplot(x="who", y="fare", col="alive",  
                    data=df, kind="bar", errorbar=None, aspect=1.0)
```

```
c:\Users\tien2\miniconda3\lib\site-packages\seaborn\axisgrid.py:118:
```

```
UserWarning: The figure layout has changed to tight
```

```
self._figure.tight_layout(*args, **kwargs)
```



1.6 1.11

```
[26]: sns.set_style("whitegrid")  
  
#---load data---  
data = pd.read_csv('data/salary.csv')  
  
#---plot the swarm plot---  
sns.swarmplot(x="gender", y="salary", data=data)  
  
ax = plt.gca()  
ax.set_title("Salary distribution")  
  
#---show plot---  
plt.show()
```



1.7 1.12

```
[27]: data = np.array([(50, 2.5), (60, 3), (65, 3.5), (70, 3.8), (75, 4), (80, 4.5),
↪(85, 5)])
data
```

```
[27]: array([[50. ,  2.5],
           [60. ,  3. ],
           [65. ,  3.5],
           [70. ,  3.8],
           [75. ,  4. ],
           [80. ,  4.5],
           [85. ,  5. ]])
```

```
[28]: X = data[:,0].reshape(-1,1)
y = data[:,1]

model = LinearRegression()
model.fit(X, y)
```

```

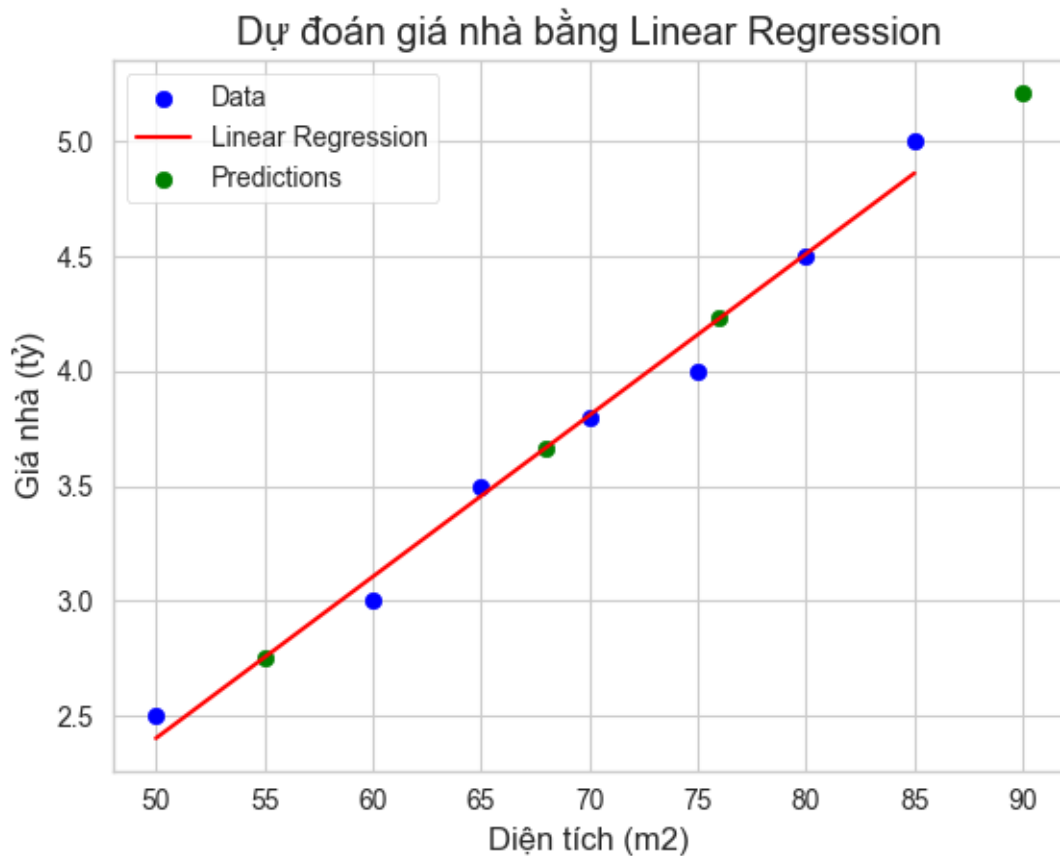
new_areas = np.array([55, 68, 76, 90]).reshape(-1, 1)
predicted_prices = model.predict(new_areas)

for area, price in zip(new_areas, predicted_prices):
    print(f"Din tích: {area} m2, Giá d đoán: {price:.2f} t")

plt.scatter(X, y, color='blue', label='Data')
plt.plot(X, model.predict(X), color='red', label='Linear Regression')
plt.scatter(new_areas, predicted_prices, color='green', label='Predictions')
plt.xlabel('Din tích (m2)')
plt.ylabel('Giá nhà (tỷ)')
plt.title('D đoán giá nhà bng Linear Regression')
plt.legend()
plt.show()

```

Din tích: [55] m2, Giá d đoán: 2.75 t
 Din tích: [68] m2, Giá d đoán: 3.67 t
 Din tích: [76] m2, Giá d đoán: 4.23 t
 Din tích: [90] m2, Giá d đoán: 5.21 t



1.8 1.13

```
[29]: heights = [[1.6], [1.65], [1.7], [1.73], [1.8]]
```

```
weights = [[60], [65], [72.3], [75], [80]]
```

```
[30]: # represents the weights of a group of people in kgs
```

```
weights = [[60], [65], [72.3], [75], [80]]
```

```
plt.title('Weights plotted against heights')
```

```
plt.xlabel('Heights in meters')
```

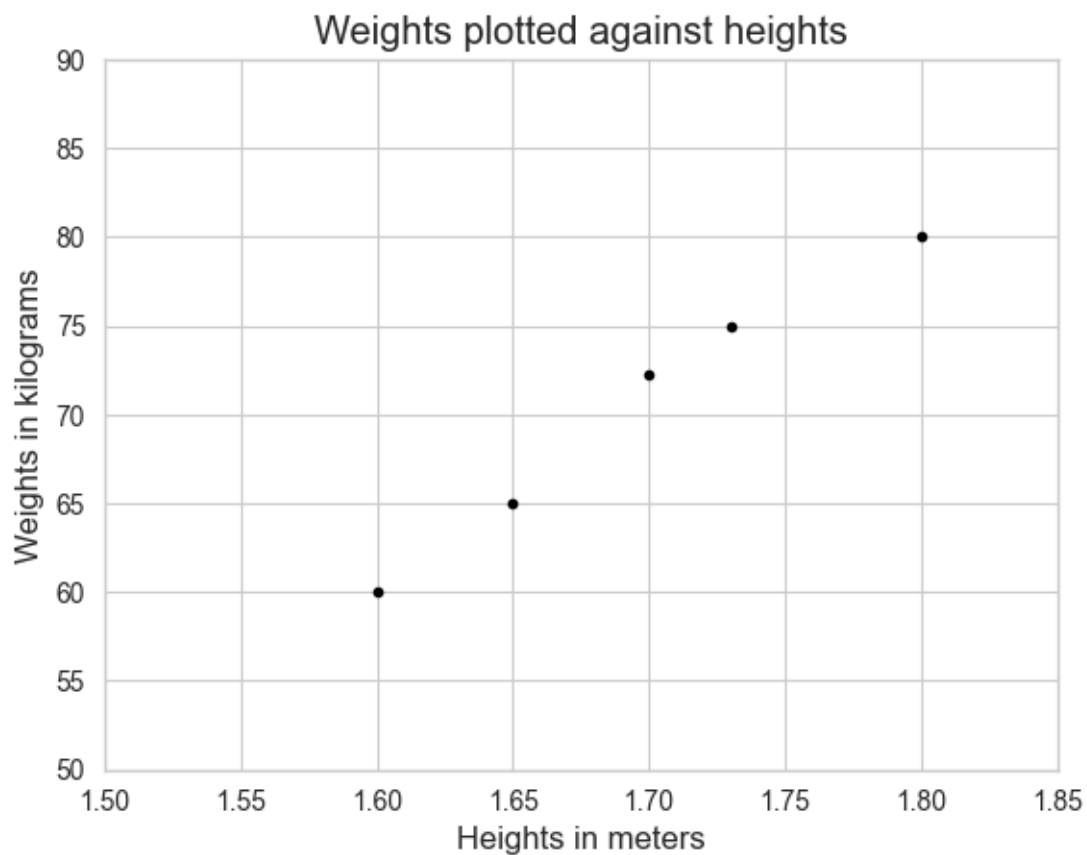
```
plt.ylabel('Weights in kilograms')
```

```
plt.plot(heights, weights, 'k.')
```

```
# axis range for x and y
```

```
plt.axis([1.5, 1.85, 50, 90])
```

```
plt.grid(True)
```



```
[31]: model = LinearRegression()
model.fit(X=heights, y=weights)

weight = model.predict([[1.75]])[0][0]
print(f'Predicted weight for height 1.75 m: {round(weight,2)} kg')
```

Predicted weight for height 1.75 m: 76.04 kg

```
[32]: import matplotlib.pyplot as plt

heights = [[1.6], [1.65], [1.7], [1.73], [1.8]]
weights = [[60], [65], [72.3], [75], [80]]

plt.title('Weights plotted against heights')
plt.xlabel('Heights in meters')
plt.ylabel('Weights in kilograms')
plt.plot(heights, weights, 'k.')

plt.axis([1.5, 1.85, 50, 90])
plt.grid(True)

# plot the regression line
plt.plot(heights, model.predict(heights), color='r')


round(model.predict([[0]])[0][0],2) # -104.75

print(round(model.intercept_[0],2)) # -104.75

print(round(model.coef_[0][0],2)) # 103.31
```

-104.75


103.31



```
[33]: plt.title('Weights plotted against heights')
plt.xlabel('Heights in meters')
plt.ylabel('Weights in kilograms')

plt.plot(heights, weights, 'k.')

plt.axis([0, 1.85, -200, 200])
plt.grid(True)
```



```
[34]: import numpy as np

print('Residual sum of squares: %.2f' %
      np.sum((weights - model.predict(heights)) ** 2))
```

Residual sum of squares: 5.34

```
[35]: # test data
heights_test = [[1.58], [1.62], [1.69], [1.76], [1.82]]
weights_test = [[58], [63], [72], [73], [85]]
```

```
[36]: # Total Sum of Squares (TSS)
weights_test_mean = np.mean(np.ravel(weights_test))
TSS = np.sum((np.ravel(weights_test) -
              weights_test_mean) ** 2)
print("TSS: %.2f" % TSS)

# Residual Sum of Squares (RSS)
RSS = np.sum((np.ravel(weights_test) -
              np.ravel(model.predict(heights_test))))
```

```

        ** 2)
print("RSS: %.2f" % RSS)

# R_squared
R_squared = 1 - (RSS / TSS)
print("R-squared: %.2f" % R_squared)

# using scikit-learn to calculate r-squared
print('R-squared: %.4f' % model.score(heights_test,
                                     weights_test))

```

TSS: 430.80
 RSS: 24.62
 R-squared: 0.94
 R-squared: 0.9429

```

[37]: import pickle

# save the model to disk
filename = './data/HeightsAndWeights_model.sav'
# write to the file using write and binary mode
pickle.dump(model, open(filename, 'wb'))

```

```

[38]: # load the model from disk
loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.score(heights_test,
                             weights_test)

result

```

[38]: 0.9428592885995254

1.8.1 Personal records

```

[39]: heights = [[1.6], [1.65], [1.7], [1.73], [1.8]]

weights = [[60], [65], [72.3], [75], [80]]

model = LinearRegression()
model.fit(heights, weights)

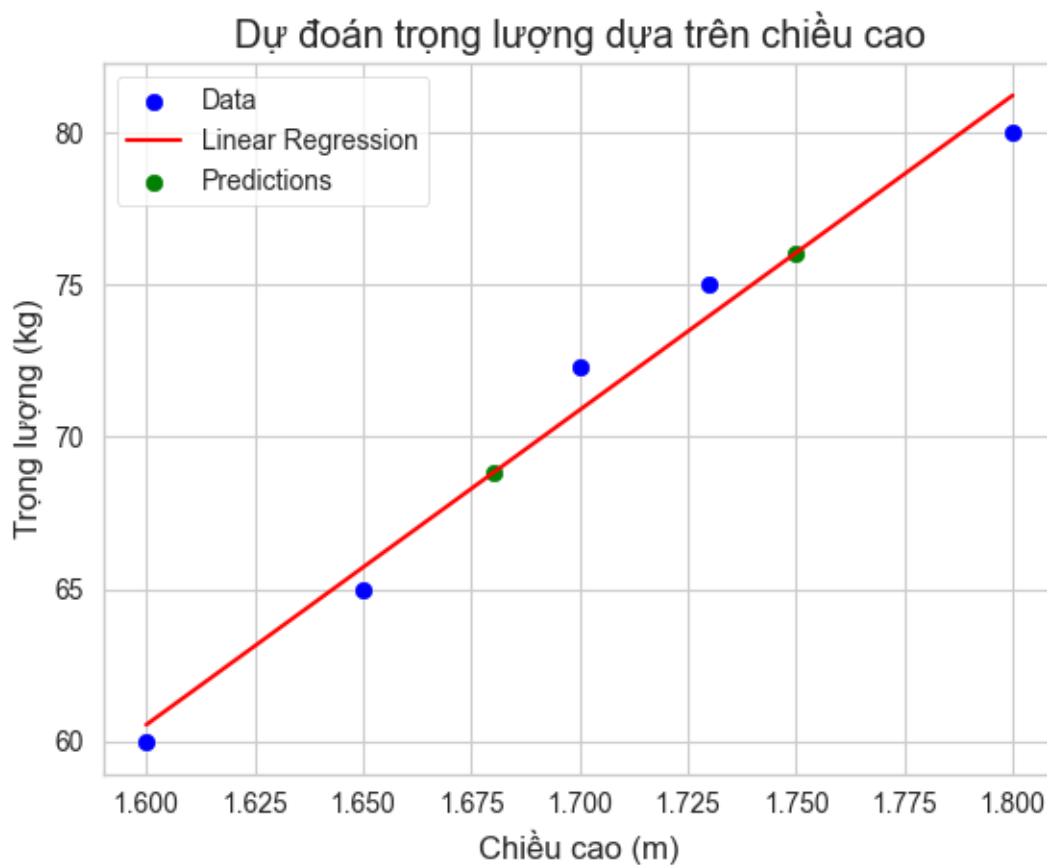
# Predict weights for new heights
new_heights = np.array([[1.68], [1.75]])
predicted_weights = model.predict(new_heights)

# Display predictions
for height, weight in zip(new_heights, predicted_weights):
    print(f"Chiu cao: {height[0]} m, Trng lng d đoán: {weight[0]:.2f} kg")

```

Chiu cao: 1.68 m, Trng lng d đoán: 68.81 kg
Chiu cao: 1.75 m, Trng lng d đoán: 76.04 kg

```
[40]: # Visualization
plt.scatter(heights, weights, color='blue', label='Data')
plt.plot(heights, model.predict(heights), color='red', label='Linear Regression')
plt.scatter(new_heights, predicted_weights, color='green', label='Predictions')
plt.xlabel('Chiu cao (m)')
plt.ylabel('Trng lng (kg)')
plt.title('D đoán trng lng da trên chiu cao')
plt.legend()
plt.show()
```



1.9 1.14

```
[41]: dataset = fetch_openml(name='boston')
dataset.data
```

c:\Users\tien2\miniconda3\lib\site-packages\sklearn\datasets_openml.py:303:
UserWarning: Multiple active versions of the dataset matching the name boston

exist. Versions may be fundamentally different, returning version 1.

```
warn(
c:\Users\tien2\miniconda3\lib\site-packages\sklearn\datasets\_openml.py:1002:
FutureWarning: The default value of `parser` will change from `liac-arff` to
`auto` in 1.4. You can set `parser='auto'` to silence this warning. Therefore,
an `ImportError` will be raised from 1.4 if the dataset is dense and pandas is
not installed. Note that the pandas parser may return different data types. See
the Notes Section in fetch_openml's API doc for details.
warn(
```

```
[41]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	
..	
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273.0	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273.0	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273.0	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273.0	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273.0	
	PTRATIO	B	LSTAT								
0	15.3	396.90	4.98								
1	17.8	396.90	9.14								
2	17.8	392.83	4.03								
3	18.7	394.63	2.94								
4	18.7	396.90	5.33								
..								
501	21.0	391.99	9.67								
502	21.0	396.90	9.08								
503	21.0	396.90	5.64								
504	21.0	393.45	6.48								
505	21.0	396.90	7.88								

[506 rows x 13 columns]

```
[42]: dataset.feature_names
```

```
[42]: ['CRIM',
      'ZN',
      'INDUS',
      'CHAS',
      'NOX',
      'RM',
      'AGE',
      'DIS',
```

```
'RAD',
'TAX',
'PTRATIO',
'B',
'LSTAT']
```

```
[43]: dataset.DESCR
```

```
[43]: """Author: \n**Source: Unknown - Date unknown \n**Please cite:
\n\nThe Boston house-price data of Harrison, D. and Rubinfeld, D.L.
'Hedonic\nprices and the demand for clean air', J. Environ. Economics &
Management,\nvvol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression
diagnostics\n...', Wiley, 1980. N.B. Various transformations are used in the
table on\npages 244-261 of the latter.\nVariables in order:\nCRIM per capita
crime rate by town\nZN proportion of residential land zoned for lots over
25,000 sq.ft.\nINDUS proportion of non-retail business acres per town\nCHAS
Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)\nNOX
nitric oxides concentration (parts per 10 million)\nRM average number of
rooms per dwelling\nAGE proportion of owner-occupied units built prior to
1940\nDIS weighted distances to five Boston employment centres\nRAD
index of accessibility to radial highways\nTAX full-value property-tax rate
per $10,000\nPTRATIO pupil-teacher ratio by town\nB 1000(Bk - 0.63)^2
where Bk is the proportion of blacks by town\nLSTAT % lower status of the
population\nMEDV Median value of owner-occupied homes in
$1000's\n\nInformation about the dataset\nCLASSTYPE: numeric\nCLASSINDEX:
last\n\nDownloaded from openml.org."
```

```
[44]: dataset.target
```

```
[44]: 0      24.0
      1      21.6
      2      34.7
      3      33.4
      4      36.2
      ...
     501     22.4
     502     20.6
     503     23.9
     504     22.0
     505     11.9
Name: MEDV, Length: 506, dtype: float64
```

```
[45]: df = pd.DataFrame(dataset.data, columns=dataset.feature_names)
      df.head()
```

```
[45]:      CRIM    ZN  INDUS  CHAS    NOX    RM    AGE    DIS  RAD    TAX  PTRATIO  \
0  0.00632  18.0   2.31    0  0.538  6.575  65.2  4.0900    1  296.0    15.3
```

1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7

	B	LSTAT
0	396.90	4.98
1	396.90	9.14
2	392.83	4.03
3	394.63	2.94
4	396.90	5.33

```
[46]: df['MEDV']=dataset.target
df.head()
```

```
[46]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

```
[47]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        506 non-null    float64
1   ZN          506 non-null    float64
2   INDUS       506 non-null    float64
3   CHAS        506 non-null    category
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         506 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    category
9   TAX         506 non-null    float64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
```

```

12 LSTAT    506 non-null    float64
13 MEDV     506 non-null    float64
dtypes: category(2), float64(12)
memory usage: 49.0 KB

```

```
[48]: print(df.isnull().sum())
```

```

CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     0
MEDV      0
dtype: int64

```

```
[49]: corr = df.corr()
corr
```

```
[49]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE \
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339
MEDV	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955

	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
CRIM	-0.379670	0.625505	0.582764	0.289946	-0.385064	0.455621	-0.388305
ZN	0.664408	-0.311948	-0.314563	-0.391679	0.175520	-0.412995	0.360445
INDUS	-0.708027	0.595129	0.720760	0.383248	-0.356977	0.603800	-0.483725
CHAS	-0.099176	-0.007368	-0.035587	-0.121515	0.048788	-0.053929	0.175260
NOX	-0.769230	0.611441	0.668023	0.188933	-0.380051	0.590879	-0.427321

RM	0.205246	-0.209847	-0.292048	-0.355501	0.128069	-0.613808	0.695360
AGE	-0.747881	0.456022	0.506456	0.261515	-0.273534	0.602339	-0.376955
DIS	1.000000	-0.494588	-0.534432	-0.232471	0.291512	-0.496996	0.249929
RAD	-0.494588	1.000000	0.910228	0.464741	-0.444413	0.488676	-0.381626
TAX	-0.534432	0.910228	1.000000	0.460853	-0.441808	0.543993	-0.468536
PTRATIO	-0.232471	0.464741	0.460853	1.000000	-0.177383	0.374044	-0.507787
B	0.291512	-0.444413	-0.441808	-0.177383	1.000000	-0.366087	0.333461
LSTAT	-0.496996	0.488676	0.543993	0.374044	-0.366087	1.000000	-0.737663
MEDV	0.249929	-0.381626	-0.468536	-0.507787	0.333461	-0.737663	1.000000


```
[50]: print(corr.abs().nlargest(3, 'MEDV').index)

print(corr.abs().nlargest(3, 'MEDV').values[:,13])
```

```
Index(['MEDV', 'LSTAT', 'RM'], dtype='object')
[1.          0.73766273  0.69535995]
```


```
[51]: plt.scatter(df['LSTAT'], df['MEDV'], marker='o')
plt.xlabel('LSTAT')
plt.ylabel('MEDV')
```

```
[51]: Text(0, 0.5, 'MEDV')
```



```
[52]: plt.scatter(df['RM'], df['MEDV'], marker='o')
plt.xlabel('RM')
plt.ylabel('MEDV')
```


```
[52]: Text(0, 0.5, 'MEDV')
```



```
[53]: fig = plt.figure(figsize=(18,15))
ax = fig.add_subplot(111, projection='3d')

ax.scatter(df['LSTAT'],
           df['RM'],
           df['MEDV'],
           c='b')

ax.set_xlabel("LSTAT")
ax.set_ylabel("RM")
ax.set_zlabel("MEDV")
plt.show()
```



```
[54]: x = pd.DataFrame(np.c_[df['LSTAT'], df['RM']], columns = ['LSTAT', 'RM'])
      Y = df['MEDV']
```

```
[55]: from sklearn.model_selection import train_test_split
      x_train, x_test, Y_train, Y_test = train_test_split(x, Y, test_size = 0.3,
      ↪random_state=5)
```

```
[56]: print(x_train.shape)
      print(Y_train.shape)
```

```
(354, 2)
```

```
(354,)
```

```
[57]: print(x_test.shape)
      print(Y_test.shape)
```

```
(152, 2)
(152,)
```

```
[58]: model = LinearRegression()
      model.fit(x_train, Y_train)
      price_prediction = model.predict(x_test)
```

```
[59]: print('R-Squared: %.4f' % model.score(x_test, Y_test))
```

```
R-Squared: 0.6162
```

```
[60]: mse = mean_squared_error(Y_test, price_prediction)
      mse
```

```
[60]: 36.49422110915324
```

```
[61]: plt.scatter(Y_test, price_prediction)
      plt.xlabel("Actual price")
      plt.ylabel("Predicted prices")
      plt.title("Actual prices vs Predicted prices")
```

```
[61]: Text(0.5, 1.0, 'Actual prices vs Predicted prices')
```



```
[62]: print(model.intercept_)
      print(model.coef_)
```

```
0.38437936780346504
[-0.65957972  4.83197581]
```

```
[63]: print(model.predict([[30,5]]))
```

```
[4.75686695]
```

```
c:\Users\tien2\miniconda3\lib\site-packages\sklearn\base.py:464: UserWarning: X
does not have valid feature names, but LinearRegression was fitted with feature
names
```

```
warnings.warn(
```

1.9.1 Plotting the 3D Hyperlane

```
[64]: import matplotlib.pyplot as plt
      import pandas as pd
      import numpy as np
      from mpl_toolkits.mplot3d import Axes3D
```

```

from sklearn.datasets import fetch_openml

dataset = fetch_openml(name='boston')

df = pd.DataFrame(dataset.data, columns=dataset.feature_names)
df['MEDV'] = dataset.target

x = pd.DataFrame(np.c_[df['LSTAT'], df['RM']], columns = ['LSTAT', 'RM'])
Y = df['MEDV']

fig = plt.figure(figsize=(18,15))
ax = fig.add_subplot(111, projection='3d')

ax.scatter(x['LSTAT'],
           x['RM'],
           Y,
           c='b')

ax.set_xlabel("LSTAT")
ax.set_ylabel("RM")
ax.set_zlabel("MEDV")

###create a meshgrid of all the values for LSTAT and RM---
x_surf = np.arange(0, 40, 1)   ###for LSTAT---
y_surf = np.arange(0, 10, 1)   ###for RM---
x_surf, y_surf = np.meshgrid(x_surf, y_surf)

from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x, Y)

###calculate z(MEDV) based on the model---
z = lambda x,y: (model.intercept_ + model.coef_[0] * x + model.coef_[1] * y)

ax.plot_surface(x_surf, y_surf, z(x_surf,y_surf),
               rstride=1,
               cstride=1,
               color='None',
               alpha = 0.4)

plt.show()

```

```

c:\Users\tien2\miniconda3\lib\site-packages\sklearn\datasets\_openml.py:303:
UserWarning: Multiple active versions of the dataset matching the name boston
exist. Versions may be fundamentally different, returning version 1.
  warn(
c:\Users\tien2\miniconda3\lib\site-packages\sklearn\datasets\_openml.py:1002:

```

FutureWarning: The default value of `parser` will change from `liac-arff` to `auto` in 1.4. You can set `parser='auto'` to silence this warning. Therefore, an `ImportError` will be raised from 1.4 if the dataset is dense and pandas is not installed. Note that the pandas parser may return different data types. See the Notes Section in `fetch_openml`'s API doc for details.

```
warn(
```

